MACHINE LEARNING FOR COMPUTATIONAL FINANCE
ASSIGNMENT 1

KARTHIK IYER

***Problem 1**: Linear Regression*

Consider a quadratic function,

$$f(x) = \frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2.$$

(1) What is the gradient ∇f and Hessian $\nabla^2 f$ of this function? Is ∇f Lipschitz continuous? If it is, what is the Lipschitz constant?

Solution: Let $\langle ., . \rangle$ denote the usual inner product in the Euclidean space. Before we proceed, let us state two facts.

(i) For any symmetric $n \times n$ matrix A, if $g(x) : \mathbb{R}^n \longrightarrow \mathbb{R}$ is defined as $g(x) = \langle Ax, x \rangle$, then $\nabla g(x) = 2Ax$. We can see why this is true by expanding the inner product and doing some algebra.

(ii) For any $m \times n$ matrix A, if $h(x) : \mathbb{R}^n \longrightarrow \mathbb{R}$ is defined as $h(x) = \langle Ax, b \rangle$ for some $b \in \mathbb{R}^m$, then $\nabla h(x) = A^T b$. We can why this is true by expanding the inner product and doing some algebra.

Now, since $f(x) : \mathbb{R}^n \longrightarrow \mathbb{R}$ is defined as $f(x) = \frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2$, we can re-write $f(x)$ as $f(x) = \frac{1}{2}(\langle \mathbf{F}x, \mathbf{F}x \rangle - 2\langle \mathbf{F}x, b \rangle + ||r||^2) = \frac{1}{2}(\langle \mathbf{F}^T\mathbf{F}x, x \rangle - 2\langle \mathbf{F}x, b \rangle + ||r||^2)$.

We can now invoke the two facts proved above to conclude that $\nabla f(x) = \mathbf{F}^T(\mathbf{F}x - b)$. (Here $\mathbf{F}$ is a $m \times n$ matrix, $x \in \mathbb{R}^n$ and $b \in \mathbb{R}^m$).

And hence, the Hessian $\nabla^2 f = \mathbf{F}^T\mathbf{F}$. It is also easy to see that ∇f is Lipschitz since

$$\begin{aligned} ||\nabla f(x) - \nabla f(y)||_2 &= ||\mathbf{F}^T\mathbf{F}(x - y)||_2 \text{ for any } x, y \in \mathbb{R}^n \\ &\leq ||\mathbf{F}^T\mathbf{F}||_2||x - y||_2 \text{ (By definition of matrix 2-norm)} \end{aligned}$$

Thus, ∇f is Lipschitz continuous with Lipschitz constant $||\mathbf{F}^T\mathbf{F}||_2 = ||\mathbf{F}||_2^2$.

(Note to self: Recall that the singular values of a $m \times n$ matrix X are the square roots of the eigenvalues of the $n \times n$ matrix X*X (where * stands for the transpose-conjugate matrix if it has complex coefficients, or the transpose if it has real coefficients). Thus, if X is $n \times n$ real symmetric matrix with non-negative eigenvalues, then eigenvalues and singular values coincide, but it is not generally the case). ■

(2) If we add a quadratic penalty to the function, $f(x) = \frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2 + \frac{\lambda}{2}||x||^2$. Answer the same set of questions in part (1) for this f.

Solution: We proceed as in part (1) to get $\nabla f(x) = \mathbf{F}^T(\mathbf{F}x - b) + \lambda x$ and $\nabla^2 f(x) = \mathbf{F}^T\mathbf{F} + \lambda I_{n \times n}$. In this case too, $\nabla f(x)$ is Lipschitz with Lipschitz constant $||\mathbf{F}||_2^2 + \lambda$. ■

(3) Implement gradient descent algorithm in **Jupyter Notebook** to solve the problem,

$$\min_x \frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2 + \frac{\lambda}{2}||x||^2.$$

***Problem 2**: LASSO*

Consider LASSO objective,

$$\min_x f(x) := \frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2 + \lambda||x||_1.$$

(1) Is f a β-smooth function? If it is, what is β? If not explain the reasons.

Solution: f is not a β- smooth function for any $\beta > 0$. Had it been so, then since $\frac{1}{2}||\mathbf{F}x - \mathbf{r}||^2$ is β smooth for $\beta = \sigma_1$, $||x||_1$ will also be $\beta^{'}$ smooth for some $\beta^{'}$. But $||x||_1$ is not C^1. ■

(2) Consider a simpler version of the problem,

$$\min_x \frac{1}{2\eta}||x - y||^2 + \lambda||x||_1$$

what is the solution (in closed form)?

Solution: We wish to minimize $\frac{1}{2\eta}||x - y||_2^2 + \lambda||x||_1$. Note that

$$\begin{aligned}
&min_x \ \frac{1}{2\eta}||y - x||_2^2 + \lambda||x||_1 \\
&= min_x \ \frac{1}{2\eta}[||y||^2 + ||x||^2 - 2\langle x, y\rangle] + \lambda||x||_1 \\
&= min_x \frac{1}{2\eta}[||x||^2 - 2\langle x, y\rangle] + \lambda||x||_1 \text{ as y is independent of x} \\
&= min_x \frac{1}{2\eta}[\sum_i x_i^2 - 2x_i y_i] + \lambda \sum_i |x_i| \\
&= min_x \frac{1}{\eta}\left[\sum_i \left(\frac{x_i^2}{2} - x_i y_i + \lambda\eta|x_i|\right)\right]. \qquad (0.1)
\end{aligned}$$

Note that we can minimize each term inside the sum separately as the terms are essentially independent from each other. Consider the problem of minimizing $L_i(x_i) = \left(\frac{x_i^2}{2} - x_i y_i + \lambda\eta|x_i|\right)$.

Note that if $y_i > 0$, then $x_i \geq 0$ (for if $x_i < 0$; then $L_i \geq 0 = L_i(0)$. Similarly, if $y_i < 0$, then $x_i \leq 0$.

If $y_i > 0$; then since $x_i \geq 0$, $L_i = \frac{x_i^2}{2} - x_i y_i + \lambda\eta x_i$. Minimizing this quantity (by setting the first derivative to 0) gives us $x_i = y_i - \lambda\eta$. Since this is feasible only when $x_i \geq 0$ and $y_i > 0$, we get $x_i = sgn(y_i)(|y_i| - \lambda\eta)^+$.

If If $y_i \leq 0$; then since $x_i \leq 0$, $L_i = \frac{x_i^2}{2} - x_i y_i - \lambda\eta x_i$. Minimizing this quantity (by setting the first derivative to 0) gives us $x_i = y_i + \lambda\eta$. Since this is feasible only when $x_i \leq 0$ and $y_i < 0$, we get $x_i = sgn(y_i)(|y_i| - \lambda\eta)^+$.

Thus in both cases we obtain $x_i = sgn(y_i)(|y_i| - \lambda\eta)^+$. Hence the desired minimizer is

$$x = (x_1, ..., x_n) \text{ where } x_i = sgn(y_i)(|y_i| - \lambda\eta)^+.$$

■

(3) Implement proximal gradient descent algorithm in **Jupyter Notebook**.

***Problem 3**: Robust Regression*

Consider Huber objective,

$$\min_x f(x) := \rho_\kappa(\mathbf{F}x - \mathbf{r}) + \lambda||x||_1$$

where ρ_κ is Huber function,

$$\rho_\kappa(a) = \sum_{i=1}^{m} \begin{cases} \kappa|a_i| - \kappa^2/2, & |a_i| > \kappa \\ a_i^2/2, & |a_i| \le \kappa \end{cases}$$

(1) For scalar case $a \in \mathbb{R}$, show that $\rho_\kappa(a) = \min_x \frac{1}{2}(x-a)^2 + \kappa|x|$. (Same derivation with Problem 2 (2)).

Proof. Let $p_\kappa(a) = \min_x \frac{1}{2}(x-a)^2 + \kappa|x|$. By Problem 2 part (2) we obtain $p_\kappa(a) = \frac{1}{2}([sgn(a)(|a| - \kappa)^+ - a]^2 + \kappa(||a| - \kappa|)^+$. For $|a| \ge \kappa$; we get

$$p_\kappa(a) = \frac{1}{2}\kappa^2 + \kappa|a - sgn(a)\kappa| = \frac{1}{2}\kappa^2 + \kappa|\,|a| - \kappa| = \kappa|a| - \frac{1}{2}\kappa^2.$$

For $|a| < \kappa$, we get $p_\kappa(a) = \frac{a^2}{2}$. This agrees with the definition of $\rho_\kappa(a)$. ■

(2) Is ρ_κ a β-smooth function? If it is what is β?

Solution: Yes. It is easy to see that ρ_κ is C^1. (As $|.|$ is smooth away from 0 and since

$$\frac{\partial \rho_\kappa}{\partial a_i}(a) = \begin{cases} \kappa sgn(a_i), & |a_i| > \kappa \\ a_i, & |a_i| \le \kappa \end{cases}$$

is continuous).

We claim that $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| \le |y_i - x_i|$.

To prove this, we assume without loss of generality that $y_i > x_i$ and consider the following 6 mutually exhaustive cases:

Case 1: $y_i \ge \kappa > x_i > -\kappa$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = |\kappa - x_i| = \kappa - x_i \le y_i - x_i = |y_i - x_i|$.

Case 2: $y_i \ge \kappa > -\kappa \ge x_i$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = \kappa + \kappa \le y_i - x_i = |y_i - x_i|$.

Case 3: $y_i > x \ge \kappa$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = 0 \le y_i - x_i = |y_i - x_i|$.

Case 4: $\kappa \ge y_i > x_i \ge -\kappa$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = y_i - x_i \le |y_i - x_i|$.

Case 5: $\kappa \ge y_i > -\kappa \ge x_i$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = y_i + \kappa \le y_i - x_i = |y_i - x_i|$.

Case 6: $-\kappa \ge y_i > x_i$. In this case, $\left|\frac{\partial \rho_\kappa}{\partial a_i}(y) - \frac{\partial \rho_\kappa}{\partial a_i}(x)\right| = 0 \le |y_i - x_i|$.

Our claim is hence justified. This claim in particular proves that ρ_κ is β smooth with $\beta = 1$. ■

(3) Is $\rho_\kappa(\mathbf{F}x - \mathbf{r})$ a β-smooth function with respect to x? If it is what is β?

Proof. Yes. Firstly, the composition of two C^1 functions is C^1. Moreover, by chain rule and the fact that $\rho_\kappa(a)$ is 1 smooth, we see that by chain rule, that for $h(x) = \rho_\kappa(\mathbf{F}x - \mathbf{r})$, $\nabla h(x) = \mathbf{F}^T \nabla \rho_\kappa(\mathbf{F}x - \mathbf{r})$. Hence

$$||\nabla h(x) - \nabla h(y)||_2 \le ||\mathbf{F}^T||_2 ||\mathbf{F}||_2 ||x - y||_2.$$

Thus $\rho_\kappa(\mathbf{F}x - \mathbf{r})$ is $\|\mathbf{F}\|_2^2$ smooth. ■

(4) Implement proximal gradient descent method in **Jupyter Notebook**.

***Problem 4**: Logistic Regression*

Assume we only care about distinguishing assets that will go up or go down, consider logistic regression objective,

$$\min_x f(x) := \sum_{i=1}^{m} \{\log(1 + \exp(\langle f^i, x\rangle)) - s^i \langle f^i, x\rangle\} + \frac{\lambda}{2}\|x\|^2$$

where $s^i \in \{-1, 1\}$, is the indicator of if the asset will go up or go down. Let's also denote $\mathbf{F} = [f^1, \dots, f^m]^T$ as we do in the note.

(1) Calculate the gradient of f, namely ∇f. Is f a β-smooth function? If it is, what is the β?

Solution: We note that $\frac{\partial f}{\partial x_j} = \sum_{i=1}^m \left(\frac{f_j^i exp\langle f^i, x\rangle}{1+exp\langle f^i, x\rangle} - s^i f_j^i \right) + \lambda x_j$ for $j = 1, 2, ..., n$.

Let $r = \mathbf{F}x$. Note that $\mathbf{F} \in \mathcal{M}_{m\times n}$ and $x \in \mathbb{R}^n$. Let $r = (r_1, r_2, ..., r_m)$. Define $exp(r) = (exp(r_1), exp(r_2), ..., exp(r_m))$.

$$\nabla f(x) = \mathbf{F}^T \left(\frac{1}{1 + exp(-r)} - s \right) + \lambda x, \tag{0.2}$$

where $\frac{1}{1+exp(r)}$ is gotten by componentwise operations.

After some algebraic manipulations, we can bound $\|\nabla^2 \mathbf{f}\| \leq \|\mathbf{F}\|_2^2 + \lambda$. Thus, f is β smooth with $\beta = \|\mathbf{F}\|_2^2 + \lambda$. (Note that this is the same β as we had for OLS regression. I am not sure if this is the optimal β though.)

■

(2) Implement gradient descent method in **Jupyter Notebook** over training data. Do cross validation to the best λ, and report the test error.

Solution: The test error turns out to be 47.2 %. ■